

CSC520 - Artificial Intelligence

Lecture 6

Dr. Scott N. Gerard

North Carolina State University

Jan 28, 2025

Agenda

- Local search
 - ▶ Hill climbing
 - ▶ Simulated annealing
 - ▶ Local beam search
 - ▶ Genetic algorithms


Local Search

- Previous search algorithms produce a *path* to the goal as a solution
 - ▶ Keep paths in memory
 - ▶ Remember alternative paths for backtracking
- In some problems, final state is solution. The path is irrelevant
 - ▶ 8-queens
 - ▶ Factory layout
 - ▶ IC design
- Local search algorithms produce the final state as the solution

Local Search and Optimization

- From a state, explore the neighboring states
- Does not keep track of paths or states that have been reached
- Uses very little memory
- Can often find reasonable solutions in large or infinite state spaces
- Aims to maximize (or minimize) an objective function defined on the state space
- Start state may not be specified


Optimization




Hill-climbing Search

```
function HILL-CLIMBING(percept) return a state  
  current  $\leftarrow$  initial_state  
  while true do  
    neighbor  $\leftarrow$  highest-valued successor state of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then  
      return current  
    current  $\leftarrow$  neighbor
```

4-Queen Problem




Start State




Goal State

4-Queen Problem


- Objective function: Number of pairwise conflicts
- Aim is to minimize the objective function
- Move one queen within its column to reduce conflicts



$h = 5$




$h = 2$



$h = 0$

Problems with Hill-climbing Search

- Local maxima: a peak that is lower than the highest peak
- Plateaus: flat area of the objective function landscape
- Ridges: sequence of local maximas



Problems with Hill-climbing Search

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	👑	13	16	13	16
👑	14	17	15	👑	14	16	16
17	👑	16	18	15	👑	15	👑
18	14	👑	15	15	14	👑	16
14	14	13	17	12	14	12	18

$$h = 17$$

- $8^8 \approx 17$ million states
- Hill-climbing solves 14% of the problem instances taking only 4 steps on average
- Gets stuck 86% of time at a local minimum
- With sideways move (same h) to escape plateau, hill-climbing solves 94% of the problem instances, but with 21 steps on average

Hill-climbing Improvements

- Stochastic hill-climbing
 - ▶ Random selection among the uphill moves
 - ▶ Selection probability proportional to the uphill steepness

Hill-climbing Improvements

- Stochastic hill-climbing
 - ▶ Random selection among the uphill moves
 - ▶ Selection probability proportional to the uphill steepness
- Random-walk hill-climbing
 - ▶ Greedy: With probability p , move to a neighbor with largest value
 - ▶ Random: With probability $1 - p$, move to a random neighbor

Hill-climbing Improvements

- Stochastic hill-climbing
 - ▶ Random selection among the uphill moves
 - ▶ Selection probability proportional to the uphill steepness
- Random-walk hill-climbing
 - ▶ Greedy: With probability p , move to a neighbor with largest value
 - ▶ Random: With probability $1 - p$, move to a random neighbor
- Random-restart hill-climbing
 - ▶ Run hill-climbing multiple times with a randomly generated initial state


Hill-climbing Improvements

- Stochastic hill-climbing
 - ▶ Random selection among the uphill moves
 - ▶ Selection probability proportional to the uphill steepness
- Random-walk hill-climbing
 - ▶ Greedy: With probability p , move to a neighbor with largest value
 - ▶ Random: With probability $1 - p$, move to a random neighbor
- Random-restart hill-climbing
 - ▶ Run hill-climbing multiple times with a randomly generated initial state
- Hill-climbing with both random-walk and random-restart

Simulated Annealing

- Inspired by *annealing* process in metallurgy used to harden metals
- A temperature variable is set to a high initial value and decreased in each iteration
- Similar to hill-climbing but instead of picking the *best* move, it picks a *random* move
- If the move improves the objective function, it is accepted
- Else accept the move with a probability that decreases with the badness of the move and with the temperature
- If temperature is decreased slowly enough, the algorithm will find global maxima

Simulated Annealing



Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a state

current \leftarrow *problem.initial_state*

for $t = 1$ to ∞ **do**

$T \leftarrow \text{schedule}(t)$

if $T = 0$ **then return** *current*

next \leftarrow random successor of *current*

$\Delta E \leftarrow \text{VALUE}(\text{current}) - \text{VALUE}(\text{next})$

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* with probability $e^{\Delta E/T}$

Local Beam Search

- Idea: Instead of keeping just 1 state, keep k best states
- Start with k randomly selected states
- Generate successors of the k states
- If any of the successors is goal, then return that successor
- Else select k best states from the successors and repeat

Local Beam Search

- Idea: Instead of keeping just 1 state, keep k best states
- Start with k randomly selected states
- Generate successors of the k states
- If any of the successors is goal, then return that successor
- Else select k best states from the successors and repeat
- Problem: all k states end up on the same local maxima
- Stochastic beam search addresses the problem by choosing k successors randomly biased toward good ones


Local Beam Search

- Idea: Instead of keeping just 1 state, keep k best states
- Start with k randomly selected states
- Generate successors of the k states
- If any of the successors is goal, then return that successor
- Else select k best states from the successors and repeat
- Problem: all k states end up on the same local maxima
- Stochastic beam search addresses the problem by choosing k successors randomly biased toward good ones
- Similar to natural selection: successors (offspring) of a state (organism) populate the next generation according to its value (fitness)

Genetic Algorithms

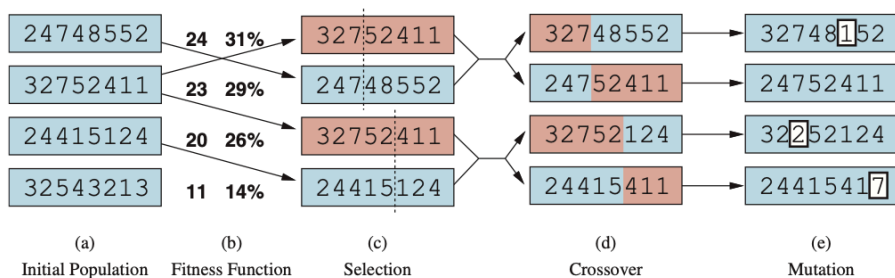
- Variant of stochastic beam search where successors are generated by combining two parents
- A state (genome) is represented as a string over a finite alphabet
- Start with k randomly generated states (population)
- An evaluation or fitness function for a state (phenome)
 - ▶ $\text{fn}(\text{genome}) \Rightarrow \text{phenome}$
- Create next generation of states by simulated evolution
 - ▶ Randomly select parents with probability proportional to their fitness
 - ▶ Combine parents using a random crossover point
 - ▶ Randomly mutate the offspring with probability equal to mutation rate

8-Queen Genetic Algorithm




- State (genome) is 8-digit string: each digit represents row number of a queen
- State = 16257483
- Phenome is a chess board with 8 queens.
- Fitness function is the number of *non – attacking* pairs
- Fitness = 27

8-Queen Genetic Algorithm



8-Queen Genetic Algorithm



Genetic Algorithms

- Random exploration can find solutions that local search cannot
- Appealing connection to human evolution (survival of the fittest)

Genetic Algorithms

- Random exploration can find solutions that local search cannot
- Appealing connection to human evolution (survival of the fittest)
- Too many tunable parameters
- Difficult to replicate performance from one problem to another
- Lack of empirical studies comparing to simpler methods
- No convincing evidence that genetic algorithms are better than hill-climbing with random restarts in general