

CSC520 - Artificial Intelligence

Lecture 7

Dr. Scott N. Gerard

North Carolina State University

Jan 30, 2025

Agenda


- Nondeterministic actions
- Partial observability
- Online search

Nondeterministic Actions

- Earlier search algorithms assumed a deterministic environment
- We now consider tasks in which actions are non-deterministic
 - ▶ Agent doesn't know what state is reached after taking an action
 - ▶ E.g. if agent executes a in s_1 , the agent may transition to s_2 , s_4 or s_5
- Instead of a sequential plan, the solution is a conditional plan (or contingency plan)

Erratic Vacuum World


- Earlier solution: [*suck*, *right*, *suck*]
- Suppose *suck* action is non-deterministic
 - ▶ On a dirty square, it cleans the square and sometimes cleans up adjacent square, too
 - ▶ On a clean square, it sometimes deposits dirt
- Generalize the transition model
 - ▶ Earlier: $\text{RESULTS}(1, \text{Suck}) = 5$
 - ▶ Now: $\text{RESULTS}(1, \text{Suck}) = \{5, 7\}$ which is a **belief state**
- Solution is a conditional plan
[*Suck*, **if** *State* = 5 **then** [*Right*, *Suck*] **else**[]]



AND-OR Search Trees

- Solutions of non-deterministic problems are *trees* rather than sequence of actions
- Agent's choices are represented as OR nodes
- Environment's choice of outcome of an action are represented as AND nodes.

Erratic Vacuum AND-OR Search Tree



[Suck, if State = 5 then [Right, Suck] else[]]


AND-OR Search Tree

- Solution is a subtree that:
 - ▶ Has a goal node at **every** leaf
 - ▶ Specifies one action at each OR node
 - ▶ Includes every outcome branch at each AND node
- Variations of BFS, DFS, A^* , etc. techniques can be used for AND-OR graph search

Slippery Vacuum World


- Right (left) action may fail leaving agent in the same location
- Keep trying the action until it succeeds
- Assumption is that each possible outcome of a non-deterministic action eventually occurs
- A cyclic solution is defined using **while** construct

[*Suck*, **while** *State* = 5 **do** *Right*, *Suck*]



Searching With No Observations

- Sensorless problem: agent's percepts provide no information
- **Belief state** is a set of physical states that an agent believes are possible
- Suppose agent knows geography but doesn't know its own location
 - ▶ Initial belief state = $\{1, 2, 3, 4, 5, 6, 7, 8\}$
 - ▶ After $[Right]$: $\{2, 4, 6, 8\}$
 - ▶ After $[Right, Suck]$: $\{4, 8\}$
 - ▶ After $[Right, Suck, Left, Suck]$: $\{7\}$
- Belief states are fully observable.
Agent always knows what it believes



Searching in Partially Observable Environments

- Earlier algorithms assumed full observability
- Searching with no observations
- Searching with partial observations


Searching With No Observations

- Search in the space of *belief states* rather than physical states
- Fully observable in belief-state space
- Solution is a sequence of actions
- Existing algorithms can be used if we formulate the search problem in terms of belief-state

Belief-state Problem


- *State space*: every possible subset of physical states; with N physical states, there can be 2^N belief states
- *Initial state*: typically a set with all physical states
- *Actions*: union of legal actions for all physical states in belief-state (Or intersection if illegal actions are possible)
- *Transition model*: belief-state resulting from applying an action in a given belief-state
$$b' = \text{RESULT}(b, a) = \{s' : s' = \text{RESULT}(s, a) \text{ and } s \in b\}$$
- *Cost function*: (tricky) cost of executing an action can be calculated from the cost of applying an action in a physical state
- *Solution* is a *path* from the start state to a goal state

Deterministic Sensorless Vacuum World



Searching In Partially Observable Environments

- Problem specification includes a $\text{PERCEPT}(s)$ function that returns percepts for a given state
- Suppose vacuum cleaner can only sense dirt its tile it (partial observability)
- Then, percept $[A, \text{Dirty}]$ can be observed in states 1 and 3, that is, the belief-state = $\{1, 3\}$
- Percept $[B, \text{Clean}]$ can be observed in states 4 and 8, that is, the belief-state = $\{4, 8\}$



Transition Model in Partially Observable Search

- Three stages of the transition model
 - ▶ Prediction stage computes belief state resulting from an action
 $\hat{b} = \text{PREDICT}(b, a)$

Transition Model in Partially Observable Search

- Three stages of the transition model

- ▶ Prediction stage computes belief state resulting from an action
 $\hat{b} = \text{PREDICT}(b, a)$
- ▶ Possible percepts stage computes the percepts that could be observed in the predicted belief state
 $\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o : o = \text{PERCEPT}(s), s \in \hat{b}\}$

Transition Model in Partially Observable Search

- Three stages of the transition model

- ▶ Prediction stage computes belief state resulting from an action
 $\hat{b} = \text{PREDICT}(b, a)$
- ▶ Possible percepts stage computes the percepts that could be observed in the predicted belief state
 $\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o : o = \text{PERCEPT}(s), s \in \hat{b}\}$
- ▶ Update stage computes the belief state that would result from the possible percepts
 $b_o = \text{UPDATE}(\hat{b}, o) = \{s : o = \text{PERCEPT}(s), s \in \hat{b}\}$

Transition Model in Partially Observable Search

- Three stages of the transition model

- ▶ Prediction stage computes belief state resulting from an action
 $\hat{b} = \text{PREDICT}(b, a)$
- ▶ Possible percepts stage computes the percepts that could be observed in the predicted belief state
 $\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o : o = \text{PERCEPT}(s), s \in \hat{b}\}$
- ▶ Update stage computes the belief state that would result from the possible percepts
 $b_o = \text{UPDATE}(\hat{b}, o) = \{s : o = \text{PERCEPT}(s), s \in \hat{b}\}$

- Combining the above stages, we get

$$\text{RESULTS}(b, a) = \{b_o : b_o = \text{UPDATE}(\text{PREDICT}(b, a), o), \\ o \in \text{POSSIBLE-PERCEPTS}(\text{PREDICT}(b, a))\}$$

Transition Model in Partially Observable Search

- Three stages of the transition model

- ▶ Prediction stage computes belief state resulting from an action
 $\hat{b} = \text{PREDICT}(b, a)$
- ▶ Possible percepts stage computes the percepts that could be observed in the predicted belief state
 $\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o : o = \text{PERCEPT}(s), s \in \hat{b}\}$
- ▶ Update stage computes the belief state that would result from the possible percepts
 $b_o = \text{UPDATE}(\hat{b}, o) = \{s : o = \text{PERCEPT}(s), s \in \hat{b}\}$


- Combining the above stages, we get

$$\text{RESULTS}(b, a) = \{b_o : b_o = \text{UPDATE}(\text{PREDICT}(b, a), o), \\ o \in \text{POSSIBLE-PERCEPTS}(\text{PREDICT}(b, a))\}$$


- PREDICT can enlarge belief state size.

UPDATE can't enlarge and may reduce belief state size.

Example




Deterministic Actions



Non-deterministic Actions

Solving Partially Observable Problems

- Use AND-OR search tree
- Solution is a conditional plan
- $[Suck, Right, \text{if } Bstate = \{6\} \text{ then } Suck \text{ else } []]$



Online Search

- We studied offline search algorithms so far
 - ▶ Compute a complete solution before taking the first action
- Online search interleaves *computation* and *action*

Online Search

- We studied offline search algorithms so far
 - ▶ Compute a complete solution before taking the first action
- Online search interleaves *computation* and *action*
- Useful for dynamic and semi-dynamic environments with penalty for delays

Online Search

- We studied offline search algorithms so far
 - ▶ Compute a complete solution before taking the first action
- Online search interleaves *computation* and *action*
- Useful for dynamic and semi-dynamic environments with penalty for delays
- Useful for non-deterministic environments since the agent can focus on contingencies that actually arise rather than those that might arise

Online Search


- We studied offline search algorithms so far
 - ▶ Compute a complete solution before taking the first action
- Online search interleaves *computation* and *action*
- Useful for dynamic and semi-dynamic environments with penalty for delays
- Useful for non-deterministic environments since the agent can focus on contingencies that actually arise rather than those that might arise
- Necessary in unknown environments
 - ▶ Agent does not know the states or effects of actions
 - ▶ Agent must learn from experience
 - ▶ E.g. a vacuum cleaner robot must learn the floor-plan of a home as it vacuums

Online Search Problem

- Agent knows the following:
 - ▶ $ACTION(s)$: the legal actions in state s
 - ▶ $c(s, a, s')$: the cost of applying action a ; agent cannot use this until it knows that s' is the outcome
 - ▶ $GOAL-TEST(s)$: the goal test
 - ▶ A heuristic function that estimates the path cost to a goal state
- Agent's objective is to reach the goal state while minimizing cost
- *Competitive ratio* is the ratio of the *actual* cost with *optimal* cost in known environment

Safely Explorable Environment

- Online search is vulnerable to *dead ends*
- No algorithm can avoid dead ends in all state spaces
- We usually assume that state space is *safely explorable*: a goal state is reachable from every state



Online Search Algorithms

- Agent can only explore the successors of the current state
 - ▶ Agent cannot bounce to a distant state in another path like A^*

Online Search Algorithms

- Agent can only explore the successors of the current state
 - ▶ Agent cannot bounce to a distant state in another path like A^*
- Depth-first search can be used assuming the actions are reversible

Online Search Algorithms


- Agent can only explore the successors of the current state
 - ▶ Agent cannot bounce to a distant state in another path like A^*
- Depth-first search can be used assuming the actions are reversible
- Can an online search agent use Hill-climbing algorithm?

Online Search Algorithms

- Agent can only explore the successors of the current state
 - ▶ Agent cannot bounce to a distant state in another path like A^*
- Depth-first search can be used assuming the actions are reversible
- Can an online search agent use Hill-climbing algorithm?
 - ▶ What about random restart Hill-climbing?

Online Search Algorithms

- Agent can only explore the successors of the current state
 - ▶ Agent cannot bounce to a distant state in another path like A^*
- Depth-first search can be used assuming the actions are reversible
- Can an online search agent use Hill-climbing algorithm?
 - ▶ What about random restart Hill-climbing?
- Random-walk hill-climbing can be used
 - ▶ Will eventually find a goal in a finite safely explorable state space
 - ▶ But can be very slow



Class Exercise

- Trace the execution of a DFS algorithm on the below simple maze problem.

