

CSC520 - Artificial Intelligence

Lecture 8

Dr. Scott N. Gerard

North Carolina State University

February 1, 2025

Multiagent Environments

Adversarial agents with conflicting goals

- Adversarial search
- MinMax algorithm
- Alpha-Beta pruning

- Games are competitive or cooperative environments with two or more agents (players)
- Different types of games
 - ▶ Deterministic vs stochastic
 - ▶ 1, 2, \dots , n players. Large $n \Rightarrow$ economy
 - ▶ Zero sum vs win-win
 - ▶ Perfect information (fully observable) vs imperfect information (partially observable)
- Study adversarial search algorithms on games
- Intent is to compute a strategy or policy which recommends a move from each state

Two-player, Zero-sum Game Problem Formulation

Deterministic, two-players, turn-taking, perfect information, zero-sum


- S_0 : initial state
- $\text{TO-MOVE}(s)$: the player to move in state s
- $\text{ACTIONS}(s)$: legal moves in state s
- $\text{RESULT}(s, a)$: transition model; state resulting from taking action a in s
- $\text{IS-TERMINAL}(s)$: if true, game is over in s
- $\text{UTILITY}(s, p)$: value of s to player p

Minimax Search

- State-space search tree
- Players alternate turns
- Compute each node's minimax value
 - ▶ Minimax value is the best achievable utility assuming both players play optimally, that is both are rational

Minimax Search Example

Two player, deterministic, zero-sum, and perfect information game



Minimax Search Algorithm

function MINIMAX-DECISION(*state*) **return** an *action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in $\text{ACTIONS}(\textit{state})$ which leads to a state with value v

function MAX-VALUE(*state*) **return** a *utility* value

if IS-TERMINAL(*state*) **return** the $\text{UTILITY}(\textit{state})$

$v \leftarrow -\infty$

for a in $\text{ACTIONS}(\textit{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\textit{state}, a)))$

return v

function MIN-VALUE(*state*) **return** a *utility* value

if IS-TERMINAL(*state*) **return** the $\text{UTILITY}(\textit{state})$


$v \leftarrow \infty$

for a in $\text{ACTIONS}(\textit{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\textit{state}, a)))$

return v

Minimax Algorithm Example



```
function MAX-VALUE(state) return a utility value
  if IS-TERMINAL(state) return the UTILITY(state)
   $v \leftarrow -\infty$ 
  for a in ACTIONS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\text{state}, a)))$ 
  return v
```

```
function MIN-VALUE(state) return a utility value
  if IS-TERMINAL(state) return the UTILITY(state)
   $v \leftarrow \infty$ 
  for a in ACTIONS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\text{state}, a)))$ 
  return v
```


Minimax Properties

- Similar to depth first search
- Time complexity: $O(b^m)$
- Space complexity: $O(bm)$


Minimax Properties

- Similar to depth first search
- Time complexity: $O(b^m)$
- Space complexity: $O(bm)$
- For example, chess has $b \approx 35, m \approx 100$
 - ▶ Too large to compute exact solution

Minimax Properties

- Similar to depth first search
- Time complexity: $O(b^m)$
- Space complexity: $O(bm)$
- For example, chess has $b \approx 35$, $m \approx 100$
 - ▶ Too large to compute exact solution
- How to reduce the complexity?
 - ▶ Dynamic pruning of the search tree
 - ▶ Early cutoff of the search tree using a heuristic evaluation function

Tree Pruning



Alpha-Beta Pruning

function ALPHA-BETA-SEARCH(*state*) **return** an *action*

$v, move \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$ **return** *move*

function MAX-VALUE(*state*, α , β) **return** (*utility*, *move*) pair

if IS-TERMINAL(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a* in ACTIONS(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{RESULT}(state, a), \alpha, \beta)$

if $v2 > v$ **then**

$v, move \leftarrow v2, a2$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $v \geq \beta$ **then return** $v, move$

return $v, move$

- α = best choice found so far for MAX (at least)

- β = best choice found so far for MIN (at most)

function MIN-VALUE(*state*, α , β) **return** (*utility*, *move*) pair

if IS-TERMINAL(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a* in ACTIONS(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{RESULT}(state, a), \alpha, \beta)$

if $v2 < v$ **then**


$v, move \leftarrow v2, a2$

$\beta \leftarrow \text{MIN}(\beta, v)$


if $v \leq \alpha$ **then return** $v, move$

return $v, move$


Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning



Alpha-Beta Pruning



Alpha-Beta Pruning




Alpha-Beta Pruning

- Pruning does not affect the final result
- Effectiveness of pruning depends upon the order of successors
 - ▶ More pruning is possible if successors likely to be best are examined first
- With *perfect ordering* time complexity drops to $O(b^{m/2})$
 - ▶ Can look ahead roughly twice as deep as minimax in the same amount of time

Class Exercise


Apply Alpha-Beta search to the game tree below.



Class Exercise


MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$

MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise


Apply Alpha-Beta search to the game tree below.



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \min(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$


MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \min(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$

MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \min(\beta, v)$



Class Exercise

MAX-VALUE prunes if $v \geq \beta$ and updates $\alpha \leftarrow \text{MAX}(\alpha, v)$

MIN-VALUE prunes if $v \leq \alpha$ and updates $\beta \leftarrow \text{MIN}(\beta, v)$

